

# ham10000

December 8, 2023

```
[ ]: %load_ext autoreload
     %autoreload 2
```

```
[ ]: import cv2
     import numpy as np
     import matplotlib.pyplot as plt
     from tqdm.notebook import tqdm
     import torch

     from torch.utils.data import DataLoader, Dataset
     from torchvision import models, transforms
     device = "cuda" if torch.cuda.is_available() else 'cpu'
     print(device)
     import wandb
     import torch.nn as nn
```

```
[ ]: wandb.login()
```

```
[ ]: from data_utils import load_dataset, LESION_TYPE
```

## 1 CLIP Zero-Shot Classification

```
[ ]: import clip
```

```
[ ]: clip_model, clip_preprocess = clip.load("ViT-B/32", device=device)
```

```
[ ]: ham_train, ham_test = load_dataset("HAM10000", transform=clip_preprocess)

     print(f"Train size: {len(ham_train)}")
     print(f"Test size: {len(ham_test)}")
     print(ham_train)
     print(ham_test)
```

```
[ ]: BATCH_SIZE = 128
```

```
[ ]: def clip_zero_shot(data_set, classes):
    # https://colab.research.google.com/drive/1IqJfogZdC61dgE4BDQILCJS-zUiphD4y?
    ↪authuser=2#scrollTo=EuZFg3ZlHOVD
    data_loader = DataLoader(data_set, batch_size=BATCH_SIZE, shuffle=True,
    ↪num_workers=2)
    # Encode text features here
    text_inputs = torch.cat([clip.tokenize(f"a photo of a {c}, a type of skin_
    ↪lesion.") for c in classes]).to(device)
    with torch.no_grad():
        text_features = clip_model.encode_text(text_inputs)
        text_features /= text_features.norm(dim=-1, keepdim=True)
    # Encode image features here
    correct = 0
    total = 0
    for image, label in tqdm(data_loader):
        image, label = image.to(device), label.to(device)
        with torch.no_grad():
            image_features = clip_model.encode_image(image)
            image_features /= image_features.norm(dim=-1, keepdim=True)
            similarity = (100.0 * image_features @ text_features.T).softmax(dim=-1)
            _, pred = similarity.max(dim=-1)
            correct += (pred == label).sum().item()
            total += len(label)

    return correct / total
```

```
[ ]: lesion_classes = LESION_TYPE.values() # This was probably only because the_
    ↪class labels were numbers, not strs
```

```
[ ]: accuracy = clip_zero_shot(data_set=ham_train, classes=lesion_classes)
    print(f"\nAccuracy = {100*accuracy:.3f}%")
```

## 2 CLIP Linear-Probe Classification

### 2.1 Logistic Regression

```
[ ]: from sklearn.linear_model import LogisticRegression
```

```
[ ]: def get_features(data_set):
    all_features = []
    all_labels = []

    with torch.no_grad():
        for images, labels in tqdm(DataLoader(data_set, batch_size=BATCH_SIZE)):
            features = clip_model.encode_image(images.to(device))
            all_features.append(features)
```

```

        all_labels.append(labels)

    return torch.cat(all_features).cpu().numpy(), torch.cat(all_labels).cpu().
    ↪numpy()

```

```

[ ]: # Calculate the image features
train_features, train_labels = get_features(ham_train)
test_features, test_labels = get_features(ham_test)

```

```

[ ]: # Perform logistic regression
classifier = LogisticRegression(random_state=0, C=0.316, max_iter=10000, ↪
    ↪verbose=1, n_jobs=-1)
classifier.fit(train_features, train_labels)

# Evaluate using the logistic regression classifier
predictions = classifier.predict(test_features)
accuracy = np.mean((test_labels == predictions).astype(float))
print(f"\nAccuracy = {100*accuracy:.3f}%")

```

## 2.2 SVM

```

[ ]: from sklearn import svm

```

```

[ ]: # Perform logistic regression
classifier = svm.SVC(random_state=0, C=0.316, max_iter=5000, verbose=1)
classifier.fit(train_features, train_labels)

# Evaluate using the logistic regression classifier
predictions = classifier.predict(test_features)
accuracy = np.mean((test_labels == predictions).astype(float))
print(f"\nAccuracy = {100*accuracy:.3f}%")

```

## 3 K-Means Clustering

```

[ ]: from scipy import stats

```

```

[ ]: def knn(x_train, y_train, x_test, y_test, K=5):
    # Needs code here
    test_pred = []
    for i in tqdm(range(len(x_test))):
        distance = np.linalg.norm(x_train - x_test[i], axis=-1)
        indices = np.argsort(distance)[:K]
        neighbors_labels = y_train[indices]
        test_pred.append(stats.mode(neighbors_labels).mode[0])

```

```

correct = (test_pred == y_test).sum()
total = len(y_test)

return correct / total

```

```
[ ]: accuracy = knn(train_features, train_labels, test_features, test_labels, K=1)
print(f"\nAccuracy = {100*accuracy:.3f}%")
```

```
[ ]: from sklearn.cluster import KMeans
```

```
[ ]: # Perform logistic regression
classifier = KMeans(n_clusters=7)
classifier.fit(train_features, train_labels)

# Evaluate using the logistic regression classifier
predictions = classifier.predict(test_features)
accuracy = np.mean((test_labels == predictions).astype(float))
print(f"\nAccuracy = {100*accuracy:.3f}%")
```

```
[ ]:
```

## 4 Random Forest

```
[ ]: from sklearn.ensemble import RandomForestClassifier
```

```
[ ]: # Perform logistic regression
classifier = RandomForestClassifier(random_state=0, verbose=1, n_jobs=-1)
classifier.fit(train_features, train_labels)

# Evaluate using the logistic regression classifier
predictions = classifier.predict(test_features)
accuracy = np.mean((test_labels == predictions).astype(float))
print(f"\nAccuracy = {100*accuracy:.3f}%")
```

## 5 ResNet 50

```
[ ]: resnet_preprocess = models.ResNet50_Weights.IMAGENET1K_V2.transforms()
weights = models.ResNet50_Weights.IMAGENET1K_V2
resnet50 = models.resnet50(weights=weights)

# Change last layer
num_features = resnet50.fc.in_features
resnet50.fc = nn.Linear(num_features, len(LESION_TYPE))

resnet50.to(device);
```

```
[ ]: from torch.optim import Adam

[ ]: train_data, test_data = load_dataset("HAM10000", transform=resnet_preprocess)

[ ]: def evaluate(model, dataloader):
    model.eval()
    with torch.no_grad():
        num_correct = 0
        total = 0
        for images, labels in tqdm(dataloader, desc="Evaluating", position=2,
↪leave=False):
            num_correct += torch.sum(labels.to(device) == torch.
↪argmax(model(images.to(device)), 1)).item()
            total += labels.size(0)
        return num_correct / total

[ ]: def train(model, optim, loss_fn, train_data, test_data, config):
    """
    Train a PyTorch model using the provided parameters.

    :param model: PyTorch model to train
    :param optim: Optimizer to use for training
    :param loss_fn: Loss function to use for training
    :param train_data: Training dataset
    :param test_data: Test dataset
    :param num_epochs: Number of epochs to train for (default is 100)
    :param batch_size: Batch size to use for data loading (default is 32)
    """
    model.train()
    run = wandb.init(
        # Set the project where this run will be logged
        project="vision-project-resnet",
        # Track hyperparameters and run metadata
        config=config)

    num_epochs = config['epochs']
    batch_size = config['batch_size']
    # Create data loaders
    train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True,
↪num_workers=2)
    test_loader = DataLoader(test_data, batch_size=batch_size, shuffle=False,
↪num_workers=2)

    for epoch in tqdm(range(num_epochs), desc="Epochs", position=0, leave=True):
        train_loss = 0.0
        correct_train = 0
        total_train = 0
```

```

    for inputs, targets in tqdm(train_loader, desc="Training", position=1,
↪leave=False):
        # Forward pass
        targets = targets.to(device)
        outputs = model(inputs.to(device))
        loss = loss_fn(outputs, targets)

        # Backward pass and optimization
        optim.zero_grad()
        loss.backward()
        optim.step()

        # Calculate train loss
        train_loss += loss.item()
        predicted = torch.argmax(outputs, 1)
        total_train += targets.size(0)
        correct_train += (predicted == targets).sum().item()

    if (epoch+1) % 2 == 0 or epoch == num_epochs - 1:
        train_loss /= len(train_loader)
        train_accuracy = correct_train / total_train

        test_accuracy = evaluate(model, test_loader)
        model.train()

        # , Test Loss: {test_loss:.4f}
        # print(f"Epoch {epoch+1}/{num_epochs}, Train Loss: {train_loss:.
↪4f}, Train Accuracy: {train_accuracy:.4f}, Test Accuracy: {test_accuracy:.
↪4f}")

        # Log metrics to wandb
        wandb.log({
            "epoch": epoch+1,
            "train_loss": train_loss,
            "train_accuracy": train_accuracy,
            "test_accuracy": test_accuracy
        })

```

```

[ ]: config = {
    "learning_rate":1e-5,
    "batch_size":64,
    "epochs":50,
    "weight_decay":1e-5,
}

```

## 5.1 Zero-Shot Resnet

```
[ ]: test_loader = DataLoader(test_data, batch_size=config['batch_size'],  
    ↪shuffle=False, num_workers=2)
```

```
[ ]: print(evaluate(resnet50, test_loader))
```

## 5.2 Fine-Tuned Resnet

```
[ ]: optim = Adam(resnet50.parameters(), lr=config['learning_rate'],  
    ↪weight_decay=config['weight_decay'])  
loss = nn.CrossEntropyLoss()
```

```
[ ]: train(resnet50, optim, loss, train_data, test_data, config)
```

```
[ ]: print(evaluate(resnet50, test_loader))
```

```
[ ]: import os  
print(os.getcwd())
```

c:\GitHub\Evaluating-CLIP-Features-for-Medical-Image-Classification

## 6 Implement a zero-shot function for medclip

```
[ ]: # implement a zero-shot function for medclip  
  
import torch  
import torchvision  
from transformers import AutoTokenizer  
from torch.utils.data import DataLoader  
from tqdm import tqdm  
  
# Device configuration  
from medclip import MedCLIPModel, MedCLIPVisionModelViT  
from medclip.modeling_medclip import MedCLIPVisionModel  
from medclip import MedCLIPProcessor  
  
# debuggin  
from PIL import Image  
  
# prepare for the demo image and texts  
from build.lib.medclip.constants import BERT_TYPE, IMG_MEAN, IMG_STD, IMG_SIZE  
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
from data_utils import load_dataset, LESION_TYPE, load_ham10000_dataset  
  
BATCH_SIZE = 64
```

```

# def medclip_zero_shot_inline(test_dataset, classes, batch_size=BATCH_SIZE):
#     # Device configuration
#     device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
#     # Data loader for the dataset
#     data_loader = DataLoader(test_dataset, batch_size=batch_size,
#                               shuffle=True, num_workers=4)
#     print(f"Device: {device}")

#     # Initialize MedClip Models
#     model = MedCLIPModel(vision_cls=MedCLIPVisionModelViT).to(device)

#     # Prepare text prompts
#     text_prompts = [f"a photo of a {c}, a type of skin lesion." for c in
#                      classes]

#     # Initialize the tokenizer
#     tokenizer = AutoTokenizer.from_pretrained(BERT_TYPE)

#     # Tokenize text prompts and convert to tensors
#     text_tokens = [tokenizer(text, return_tensors='pt', padding=True,
#                               truncation=True, add_special_tokens=True) for text in text_prompts]

#     # Encode text prompts using MedClip's text model
#     # Inside the medclip_zero_shot function
#     text_features = [
#         model.encode_text(
#             input_ids=tokens['input_ids'].to(device),
#             attention_mask=tokens['attention_mask'].to(device)
#         )
#         for tokens in text_tokens
#     ]

#     # Initialize variables for accuracy calculation
#     correct = 0
#     total = 0

#     for images, labels in tqdm(data_loader):
#         images, labels = images.to(device), labels.to(device)

#         # TODO: Encode images using MedClip's vision model
#         image_features = model.encode_image(images)

#         # Flatten text_features into a single 2D tensor
#         text_features_tensor = torch.cat(text_features, dim=0)

#         # Calculate similarity and make predictions
#         similarity = torch.matmul(image_features, text_features_tensor.t())
#         _, predictions = similarity.max(dim=-1)

```



```

#         # Update correct and total counts
#         correct += (predictions == labels).sum().item()
#         total += labels.size(0)

#     return correct / total

# # Load HAM10000 dataset
# transform = torchvision.transforms.Compose([
#     torchvision.transforms.Resize((IMG_SIZE, IMG_SIZE)),
#     torchvision.transforms.ToTensor(),
#     torchvision.transforms.Normalize(mean=[IMG_MEAN], std=[IMG_STD])
# ])

# # train_dataset, test_dataset = load_ham10000_dataset(transform=transform,
# ↪data_dir='data/ham10000')
# train_dataset, test_dataset = load_dataset("HAM10000", transform=transform,
# ↪data_dir='data/ham10000/')
# classes = list(LESION_TYPE.values()) # From the data_utils.py file

# # Run zero-shot classification
# acc = medclip_zero_shot_inline(test_dataset, classes)
# print(f"Accuracy: {acc:.2f}")

```

```

[ ]: def medclip_zero_shot(model, test_dataset, classes, batch_size=BATCH_SIZE):
    # Data loader for the dataset
    data_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=True,
    ↪num_workers=4)

    # Prepare text prompts
    text_prompts = [f"a photo of a {c}, a type of skin lesion." for c in
    ↪classes]

    # Initialize the tokenizer
    tokenizer = AutoTokenizer.from_pretrained(BERT_TYPE)
    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Device: {device}")

    # Tokenize text prompts and convert to tensors
    text_tokens = [tokenizer(text, return_tensors='pt', padding=True,
    ↪truncation=False, add_special_tokens=True) for text in text_prompts]

    # Encode text prompts using MedClip's text model
    # Inside the medclip_zero_shot function
    text_features = [
        model.encode_text(
            input_ids=tokens['input_ids'].to(device),
            attention_mask=tokens['attention_mask'].to(device)

```

```

    )
    for tokens in text_tokens
]

# Initialize variables for accuracy calculation
correct = 0
total = 0

for images, labels in tqdm(data_loader):
    images, labels = images.to(device), labels.to(device)
    # Encode images using MedClip's vision model
    # with torch.no_grad():
    image_features = model.encode_image(images)
    # Flatten text_features into a single 2D tensor
    text_features_tensor = torch.cat(text_features, dim=0)

    # Calculate similarity and make predictions
    similarity = torch.matmul(image_features, text_features_tensor.t())
    _, predictions = similarity.max(dim=-1)

    # Update correct and total counts
    correct += (predictions == labels).sum().item()
    total += len(labels)

return correct / total

```

## 6.1 Load HAM10000 dataset and test MedClip's zero-shot capabilities

```

[ ]: transform = torchvision.transforms.Compose([
    torchvision.transforms.Resize((IMG_SIZE, IMG_SIZE)),
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize(mean=[IMG_MEAN], std=[IMG_STD])
])

ham_train, ham_test = load_ham10000_dataset(data_dir="data/ham10000/",
    ↪transform=transform)
classes = list(LESION_TYPE.values()) # From the data_utils.py file

```

Loading HAM10000 dataset...

MedCLIP\_ResNet50\_model

```

[ ]: # load MedCLIP-ResNet50
MedCLIP_ResNet50_model = MedCLIPModel(vision_cls=MedCLIPVisionModel).to(device)
accuracy = medclip_zero_shot(MedCLIP_ResNet50_model, ham_train, classes)
print(f"\nAccuracy = {100*accuracy:.3f}%")

```

Some weights of the model checkpoint at emilyalsentzer/Bio\_ClinicalBERT were not

used when initializing BertModel: ['cls.predictions.transform.dense.bias', 'cls.seq\_relationship.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias', 'cls.predictions.transform.dense.weight', 'cls.seq\_relationship.weight', 'cls.predictions.bias', 'cls.predictions.decoder.weight']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Device: cuda

100%| | 141/141 [00:23<00:00, 6.11it/s]

Accuracy = 22.346%

MedCLIP\_ViT\_model

```
[ ]: # load MedCLIP-ViT
MedCLIP_ViT_model = MedCLIPModel(vision_cls=MedCLIPVisionModelViT).to(device)
accuracy = medclip_zero_shot(MedCLIP_ViT_model, ham_train, classes)
print(f"\nAccuracy = {100*accuracy:.3f}%")
```

Some weights of the model checkpoint at microsoft/swin-tiny-patch4-window7-224 were not used when initializing SwinModel: ['classifier.bias',

'classifier.weight']

- This IS expected if you are initializing SwinModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing SwinModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of the model checkpoint at emilyalsentzer/Bio\_ClinicalBERT were not used when initializing BertModel: ['cls.predictions.transform.dense.bias',

'cls.seq\_relationship.bias', 'cls.predictions.transform.LayerNorm.weight',

'cls.predictions.transform.LayerNorm.bias',

'cls.predictions.transform.dense.weight', 'cls.seq\_relationship.weight',

'cls.predictions.bias', 'cls.predictions.decoder.weight']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Device: cuda

100%| | 141/141 [00:25<00:00, 5.63it/s]

Accuracy = 27.593%

## 6.2 Load NIH Chest X-ray dataset

```
[ ]: import os
      # os.chdir('../')
      print(os.getcwd())
```

c:\GitHub\Evaluating-CLIP-Features-for-Medical-Image-Classification

```
[ ]: import torch
      import torchvision
      import torch.nn.functional as F
      from tqdm import tqdm
      from transformers import AutoTokenizer
      from torch.utils.data import DataLoader

      # Device configuration
      from data_utils import load_nih_dataset_split, NIH_CLASS_TYPES, load_dataset
      from medclip import MedCLIPModel, MedCLIPVisionModelViT, MedCLIPVisionModel
      from build.lib.medclip.constants import BERT_TYPE, IMG_MEAN, IMG_STD, IMG_SIZE

      # debuggin
      from PIL import Image

      BATCH_SIZE = 128

      transform = torchvision.transforms.Compose([
          torchvision.transforms.Resize((IMG_SIZE, IMG_SIZE)),
          torchvision.transforms.ToTensor(),
          torchvision.transforms.Normalize(mean=[IMG_MEAN], std=[IMG_STD])
      ])

      device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
      # NIH_CLASS_TYPES
      classes = list(NIH_CLASS_TYPES) # From the data_utils.py file
      classes

      # nih_train, nih_test = load_nih_dataset_split(transform=transform)
      nih_train, nih_test = load_dataset("NIH", transform=transform, data_dir='data/
      ↪nih/')
```

NIH Dataset: Compose(  
 Resize(size=(224, 224), interpolation=bilinear, max\_size=None,

```

antialias=warn)
    ToTensor()
    Normalize(mean=[0.5862785803043838], std=[0.27950088968644304])
)

```

```

[ ]: def medclip_zero_shot(model, test_dataset, classes, batch_size=BATCH_SIZE):
    # Data loader for the dataset
    data_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=True,
    ↪ num_workers=2)

    # Prepare text prompts
    text_prompts = [f"a photo of a {c}, a type of Chest x ray." for c in
    ↪ classes]

    # Initialize the tokenizer
    tokenizer = AutoTokenizer.from_pretrained(BERT_TYPE)

    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    print(f"Device: {device}")

    # Tokenize text prompts and convert to tensors
    text_tokens = [tokenizer(text, return_tensors='pt', padding=True,
    ↪ truncation=False, add_special_tokens=True) for text in text_prompts]
    # print('text_tokens', text_prompts)
    # Encode text prompts using MedClip's text model
    # Inside the medclip_zero_shot function
    text_features = [
        model.encode_text(
            input_ids=tokens['input_ids'].to(device),
            attention_mask=tokens['attention_mask'].to(device)
        )
        for tokens in text_tokens
    ]

    # Initialize variables for accuracy calculation
    correct = 0
    total = 0
    # print('text_features', text_features)
    for images, labels in tqdm(data_loader):
        images, labels = images.to(device), labels.to(device)
        # Encode images using MedClip's vision model
        # with torch.no_grad():
        image_features = model.encode_image(images)
        # Flatten text_features into a single 2D tensor
        text_features_tensor = torch.cat(text_features, dim=0)

        # Calculate similarity and make predictions
        similarity = torch.matmul(image_features, text_features_tensor.t())

```

```

_, predictions = similarity.max(dim=-1)

# Update correct and total counts
correct += (predictions == labels).sum().item()
total += len(labels)

return correct / total

```

Load MedCLIP-ResNet50

```

[ ]: MedCLIP_ResNet50_model = MedCLIPModel(vision_cls=MedCLIPVisionModel).to(device)
MedCLIP_ResNet50_model
accuracy = medclip_zero_shot(MedCLIP_ResNet50_model, nih_train, classes)
print(f"\nAccuracy = {100*accuracy:.3f}%")

```

c:\Users\mario\anaconda3\Lib\site-packages\torchvision\models\\_utils.py:208:

UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.

warnings.warn(

c:\Users\mario\anaconda3\Lib\site-packages\torchvision\models\\_utils.py:223:

UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=ResNet50\_Weights.IMAGENET1K\_V1`. You can also use `weights=ResNet50\_Weights.DEFAULT` to get the most up-to-date weights.

warnings.warn(msg)

Some weights of the model checkpoint at emilyalsentzer/Bio\_ClinicalBERT were not used when initializing BertModel: ['cls.predictions.transform.dense.bias',

'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.LayerNorm.bias', 'cls.predictions.decoder.weight', 'cls.predictions.transform.dense.weight', 'cls.seq\_relationship.bias', 'cls.seq\_relationship.weight', 'cls.predictions.bias']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Device: cuda

100% | 789/789 [04:57<00:00, 2.66it/s]

Accuracy = 53.138%

Load MedCLIP-ViT

```
[ ]: MedCLIP_ViT_model = MedCLIPModel(vision_cls=MedCLIPVisionModelViT).to(device)
accuracy = medclip_zero_shot(MedCLIP_ViT_model, nih_train, classes)
print(f"\nAccuracy = {100*accuracy:.3f}%")
```

c:\Users\mario\anaconda3\Lib\site-packages\torch\functional.py:504: UserWarning: torch.meshgrid: in an upcoming release, it will be required to pass the indexing argument. (Triggered internally at

C:\cb\pytorch\_1000000000000\work\aten\src\ATen\native\TensorShape.cpp:3527.)

```
    return _VF.meshgrid(tensors, **kwargs) # type: ignore[attr-defined]
Some weights of the model checkpoint at microsoft/swin-tiny-patch4-window7-224
were not used when initializing SwinModel: ['classifier.weight',
'classifier.bias']
```

- This IS expected if you are initializing SwinModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing SwinModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model). Some weights of the model checkpoint at emilyalsentzer/Bio\_ClinicalBERT were not used when initializing BertModel: ['cls.predictions.transform.LayerNorm.bias', 'cls.seq\_relationship.bias', 'cls.predictions.transform.dense.bias', 'cls.predictions.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.seq\_relationship.weight', 'cls.predictions.transform.dense.weight', 'cls.predictions.decoder.weight']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Device: cuda

100%| | 789/789 [3:36:49<00:00, 16.49s/it]

Accuracy = 16.531%

```
[ ]: import numpy as np
def get_features(data_set, model):
    all_features = []
    all_labels = []

    with torch.no_grad():
        for images, labels in tqdm(DataLoader(data_set, batch_size=BATCH_SIZE)):
            features = model.encode_image(images.to(device))
            all_features.append(features)
```

```

        all_labels.append(labels)

    return torch.cat(all_features).cpu().numpy(), torch.cat(all_labels).cpu().
    ↪numpy()

```

```

[ ]: MedCLIP_ResNet50_model = MedCLIPModel(vision_cls=MedCLIPVisionModel).to(device)

# Calculate the image features
train_features, train_labels = get_features(nih_train, MedCLIP_ResNet50_model)
test_features, test_labels = get_features(nih_test, MedCLIP_ResNet50_model)

```

c:\Users\mario\anaconda3\Lib\site-packages\torchvision\models\\_utils.py:208:

UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.

```
warnings.warn(
```

c:\Users\mario\anaconda3\Lib\site-packages\torchvision\models\\_utils.py:223:

UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=ResNet50\_Weights.IMAGENET1K\_V1`. You can also use `weights=ResNet50\_Weights.DEFAULT` to get the most up-to-date weights.

```
warnings.warn(msg)
```

Some weights of the model checkpoint at emilyalsentzer/Bio\_ClinicalBERT were not used when initializing BertModel: ['cls.predictions.transform.LayerNorm.bias',

```

'cls.seq_relationship.bias', 'cls.predictions.transform.dense.bias',
'cls.predictions.bias', 'cls.predictions.transform.LayerNorm.weight',
'cls.seq_relationship.weight', 'cls.predictions.transform.dense.weight',
'cls.predictions.decoder.weight']

```

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

```
100%|      | 789/789 [20:23<00:00, 1.55s/it]
```

```
100%|      | 88/88 [02:14<00:00, 1.52s/it]
```

```

[ ]: from sklearn.linear_model import LogisticRegression
# Perform logistic regression
classifier = LogisticRegression(random_state=0, C=0.316, max_iter=10000,
    ↪verbose=1, n_jobs=-1)
classifier.fit(train_features, train_labels)

# Evaluate using the logistic regression classifier
predictions = classifier.predict(test_features)
accuracy = np.mean((test_labels == predictions).astype(float))
print(f"\n MedClip ResNet50 NIH Image Features Accuracy = {100*accuracy:.3f}%")

```



```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 16 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 1.5min finished
```

MedClip ResNet50 NIH Image Features Accuracy = 54.995%

```
[ ]: # same thing for ViT  
MedCLIP_ViT_model = MedCLIPModel(vision_cls=MedCLIPVisionModelViT).to(device)  
  
# Calculate the image features  
train_features, train_labels = get_features(nih_train, MedCLIP_ViT_model)  
test_features, test_labels = get_features(nih_test, MedCLIP_ViT_model)
```

Some weights of the model checkpoint at microsoft/swin-tiny-patch4-window7-224 were not used when initializing SwinModel: ['classifier.weight', 'classifier.bias']

- This IS expected if you are initializing SwinModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing SwinModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of the model checkpoint at emilyalsentzer/Bio\_ClinicalBERT were not used when initializing BertModel: ['cls.predictions.transform.LayerNorm.bias', 'cls.seq\_relationship.bias', 'cls.predictions.transform.dense.bias', 'cls.predictions.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.seq\_relationship.weight', 'cls.predictions.transform.dense.weight', 'cls.predictions.decoder.weight']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

100%| | 789/789 [26:12<00:00, 1.99s/it]

100%| | 88/88 [02:53<00:00, 1.97s/it]

```
[ ]: from sklearn.linear_model import LogisticRegression  
# Perform logistic regression  
classifier = LogisticRegression(random_state=0, C=0.316, max_iter=10000,   
    ↪ verbose=1, n_jobs=-1)  
classifier.fit(train_features, train_labels)  
  
# Evaluate using the logistic regression classifier  
predictions = classifier.predict(test_features)  
accuracy = np.mean((test_labels == predictions).astype(float))  
print(f"\n MedClip ViT NIH Image Features Accuracy = {100*accuracy:.3f}%")
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 16 concurrent workers.  
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 1.6min finished
```

MedClip ViT NIH Image Features Accuracy = 55.342%

```
[ ]: # now for HAM10000  
MedCLIP_ResNet50_model = MedCLIPModel(vision_cls=MedCLIPVisionModel).to(device)  
  
# Calculate the image features  
train_features, train_labels = get_features(ham_train, MedCLIP_ResNet50_model)  
test_features, test_labels = get_features(ham_test, MedCLIP_ResNet50_model)
```

```
c:\Users\mario\anaconda3\Lib\site-packages\torchvision\models\_utils.py:208:  
UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be  
removed in the future, please use 'weights' instead.  
warnings.warn(  
c:\Users\mario\anaconda3\Lib\site-packages\torchvision\models\_utils.py:223:  
UserWarning: Arguments other than a weight enum or `None` for 'weights' are  
deprecated since 0.13 and may be removed in the future. The current behavior is  
equivalent to passing `weights=ResNet50_Weights.IMAGENET1K_V1`. You can also use  
`weights=ResNet50_Weights.DEFAULT` to get the most up-to-date weights.  
warnings.warn(msg)  
Some weights of the model checkpoint at emilyalsentzer/Bio_ClinicalBERT were not  
used when initializing BertModel: ['cls.predictions.transform.LayerNorm.bias',  
'cls.seq_relationship.bias', 'cls.predictions.transform.dense.bias',  
'cls.predictions.bias', 'cls.predictions.transform.LayerNorm.weight',  
'cls.seq_relationship.weight', 'cls.predictions.transform.dense.weight',  
'cls.predictions.decoder.weight']  
- This IS expected if you are initializing BertModel from the checkpoint of a  
model trained on another task or with another architecture (e.g. initializing a  
BertForSequenceClassification model from a BertForPreTraining model).  
- This IS NOT expected if you are initializing BertModel from the checkpoint of  
a model that you expect to be exactly identical (initializing a  
BertForSequenceClassification model from a BertForSequenceClassification model).  
100%|      | 141/141 [01:31<00:00, 1.54it/s]  
100%|      | 16/16 [00:06<00:00, 2.33it/s]
```

```
[ ]: from sklearn.linear_model import LogisticRegression  
# Perform logistic regression  
classifier = LogisticRegression(random_state=0, C=0.316, max_iter=10000,  
    verbose=1, n_jobs=-1)  
classifier.fit(train_features, train_labels)  
  
# Evaluate using the logistic regression classifier  
predictions = classifier.predict(test_features)  
accuracy = np.mean((test_labels == predictions).astype(float))
```

```
print(f"\n MedClip ResNet50 HAM1000 Image Features Accuracy = {100*accuracy:.3f}%")
```

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 16 concurrent workers.
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 2.4s finished
```

MedClip ResNet50 HAM1000 Image Features Accuracy = 73.054%

```
[ ]: # same thing for ViT
MedCLIP_ViT_model = MedCLIPModel(vision_cls=MedCLIPVisionModelViT).to(device)

# Calculate the image features
train_features, train_labels = get_features(ham_train, MedCLIP_ViT_model)
test_features, test_labels = get_features(ham_test, MedCLIP_ViT_model)

from sklearn.linear_model import LogisticRegression
# Perform logistic regression
classifier = LogisticRegression(random_state=0, C=0.316, max_iter=10000,
    verbose=1, n_jobs=-1)
classifier.fit(train_features, train_labels)

# Evaluate using the logistic regression classifier
predictions = classifier.predict(test_features)
accuracy = np.mean((test_labels == predictions).astype(float))
print(f"\n MedClip ViT HAM1000 Image Features Accuracy = {100*accuracy:.3f}%")
```

Some weights of the model checkpoint at microsoft/swin-tiny-patch4-window7-224 were not used when initializing SwinModel: ['classifier.weight', 'classifier.bias']

- This IS expected if you are initializing SwinModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing SwinModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some weights of the model checkpoint at emilyalsentzer/Bio\_ClinicalBERT were not used when initializing BertModel: ['cls.predictions.transform.LayerNorm.bias',

'cls.seq\_relationship.bias', 'cls.predictions.transform.dense.bias', 'cls.predictions.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.seq\_relationship.weight', 'cls.predictions.transform.dense.weight', 'cls.predictions.decoder.weight']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

```
100%|      | 141/141 [04:20<00:00, 1.85s/it]
100%|      | 16/16 [00:19<00:00, 1.24s/it]
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 16 concurrent workers.
[Parallel(n_jobs=-1)]: Done 1 out of 1 | elapsed: 1.7s finished
```

MedClip ViT HAM1000 Image Features Accuracy = 74.152%